

A Study on Soft Core Processor Configurations for Embedded Applications

Benjamin Kueffler¹, Rafi Meguerdijian², Christopher Cole³, Vlad Cretu⁴, Sasoun Torousian⁵, and Mohamed El-Hadedy⁶

¹ Department of Electrical and Computer Engineering, California Polytechnic State University, Pomona - bkkueffler@cpp.edu

² Department of Electrical and Computer Engineering, California Polytechnic State University, Pomona - meguerdijian@cpp.edu

³ Department of Electrical and Computer Engineering, California Polytechnic State University - crcole@cpp.edu

⁴ Department of Electrical and Computer Engineering, California Polytechnic State University - vfcretu@cpp.edu

⁵ Department of Electrical and Computer Engineering, California Polytechnic State University - storousian@cpp.edu

⁶ Department of Electrical and Computer Engineering, California Polytechnic State University - mealy@cpp.edu

Abstract

Field Programmable Gate Array (FPGA) designs have become widely adopted with their flexibility in accelerating performance of specific applications at a low power cost. This flexibility comes with the desire to also have customizable designs to suit the application. In this survey, we will perform tests on the Xilinx MicroBlaze processor within a 7-Series FPGA. Each of these tests enables some features on the processor while disabling others. Minimum area, high performance, maximum frequency, and frequency optimized implementations will be produced in order to survey tradeoffs that exist within these various architectures. The application that will be chosen to survey these tradeoffs will be an encryption and decryption algorithm. Our tests will demonstrate the tradeoff of power and area for performance, with an emphasis on tradeoffs for embedded cryptography applications.

Index Terms

FPGA, hardware/software interface, MicroBlaze, soft-core processor, cryptography

I. INTRODUCTION

SINCE their introduction in 1985, Field Programmable Gate Array (FPGA) designs have gained popularity as a tool to create digital circuits without the cost and complexities associated with application specific integrated circuits (ASIC). Their reprogrammability have provided designers with a way to achieve hardware acceleration while still being able to rapidly change the design. Overhead costs associated with programmable logic kept flexible solutions such as processor cores out of FPGA designs. However, since their inception, FPGAs have seen capacities increase ten thousand times, and have seen performance increase one-hundred fold [1]. This has given architects the ability to implement processors within the FPGA fabric, to provide additional flexibility and computational variety to their designs.

The importance of software alongside hardware in systems is substantial, as the ability for software configuration can optimize the performance driven hardware on the chip in real time. In addition to flexibility, software partitioning can be used to drive power costs of chips down on the order of 90%, by only enabling the highest performing clusters of hardware at a time [2]. Beyond this, the hardware/software interface is useful in providing additional debugging and real time updates by reading the status of particular hardware sectors.

One of the main competitors of soft processors are their hard counterparts. System designers have often chosen an off chip processing solution to incorporate alongside their separate FPGA device. However, in recent years, system on a chip

(SoC) devices have grown in popularity, allowing for both programmable logic and hard processor to be present within the same fabric. The Xilinx Zynq family is an example of this, with a single or dual core ARM processor at the disposal of the software designer to communicate directly with the FPGA over the Advanced eXtensible Interface (AXI) bus [3]. These hard processors give a challenge for their soft counterparts by offering faster speeds and higher performance per Watt [4]. However, soft processors offer up the potential for reduced cost and greater flexibility, with the ability to change processor architectures or features via reprogramming.

The MicroBlaze is a 32/64 bit customizable soft processor that can be used alongside any Xilinx FPGA based system. This soft processor contains the configurability needed to reduce cost to make it preferable to hard processors in many applications. The MicroBlaze is able to be customized in 32 or 64 bit modes, allowing for both large and small data widths to be implemented. In recent years, a floating point unit was added as a customizable option, allowing for more taxing floating point programs to be run on the core. More granularity exists within the microblaze system by allowing optional barrel shifters, multipliers, and divider units to be generated through the GUI. This can provide the option for small embedded applications to forgo these features in order to achieve a very small processor. Such minimal processors can achieve utilization as low as 600 LUTs and 300 registers in Spartan 7 series FPGA devices [5]. More intensive applications may choose to utilize these features, achieving higher performance.

The system architecture consists of a typical FPGA design utilizing a soft processor such as the MicroBlaze. The MicroBlaze system, along with many other designs, utilizes AMBA AXI as the bus. The processor acts as a master in order to control peripherals and access memory. The data and instruction memory lies on block RAM acting as a slave on the AXI bus. Additional slaves utilized for the purpose of testing include the AXI Uartlite and AXI timer submodules. The uartlite acts as an interface between AXI and UART and allows tests to occur by loading in files to encrypt, and reading out the result for verification. The AXI timer is enabled by the processor at the start of the algorithm under test, and is critical for giving cycle-accurate readings of the execution time. A clock generator instantiates a phase locked loop (PLL) in order to produce the system's clock, which is set to the maximum achievable frequency in the particular implementation.

IV. IMPLEMENTATIONS

Six distinct implementations of the processor were chosen in order to achieve a variety of tradeoffs for analysis. Note that in each of these implementations, regardless of target, the debug mode feature was included in the MicroBlaze, this increases area by a constant amount among all the tests (about 500 LUTs/Registers).

A. Three Stage Pipeline

This build targets the minimum area achievable. It utilizes a three stage pipeline. This type of processor build is frequently used among FPGA designers working in the embedded field, because it allows easy entry into building a software system with minimum sacrifice of chip space. This build sacrifices performance by giving up additional hardware features, but will remain cheap in terms of power and area.

B. Three Stage Pipeline with Multiplier and Divider

This build augments the normal three stage pipeline by adding a multiplier and divider within the execution stage. This uses up additional space on the chip, but also decreases the maximum achievable frequency to 115 MHz.

C. Five Stage Pipeline

The five stage pipeline represents the most used MicroBlaze pipeline. It is the pipelined used for the maximum performance of the MicroBlaze. In this implementation, we use the five stage pipeline with no multiplier, divider, or branch target buffer (BTB). This barebones pipeline provides a baseline for the performance of the five stage pipeline.

D. Five Stage Pipeline with Multiplier

This implementation augments the five stage pipeline to include a multiplier unit in the execution stage. This reduces the execution time of the program significantly. Due to the DSP slice resource on Xilinx FPGAs, the multiplier does not add higher resource utilization of the LUTs or flip flops.

E. Five Stage Pipeline with Multiplier and Divider

This build adds a multiplier and a divider unit in the execution stage of the five stage pipeline. This greatly reduces the execution time of the program. Resource utilization goes up a moderate amount from the addition of the divider.

F. Five Stage Pipeline with Multiplier, Divider, and Branch Target Buffer

Uses the same five stage pipeline with the multiplier and divider, but adds an additional branch target buffer. The BTB enhances the branch prediction capabilities of the microblaze by adding an additional cache to store the results of previous branches. In a vacuum, this would increase performance for most designs, but the MicroBlaze implementation of the BTB does reduce the maximum frequency by about 20%. This means that to see a performance increase, the addition of the BTB must make up for the lost maximum frequency.

V. SOFTWARE AND TEST

In order to give an accurate depiction of the tradeoffs of the MicroBlaze system, a software application was produced that could take advantage of additional hardware acceleration. Encrypting and decrypting messages takes significant computational resources, and can be scaled with better hardware. The RSA algorithm was used, which uses prime numbers as a basis for encoding and decoding. One public key is generated for encryption, while one private key is used for decrypting the encrypted message.

Both encryption and decryption require modulo and multiplication to encode and decode. These operations can be accelerated with multiply and divide operations, which require additional hardware to be present in processors. Without this hardware, the performance of the encoding and decoding will be reduced. This tradeoff of area and power in return for performance was the basis for the tests done using this algorithm.

The test consisted of encrypting and decrypting two scripts. One of the scripts consisted of 32kB of latin text, and the other script consisted of 32kB of numbers. Different scripts were used in order to verify that the algorithm was working as intended with all characters and to record any potential variations in performance with a variety of characters. The test data was large in order to give a more accurate result. All the data from the scripts was loaded in over UART into the internal block RAM of the FPGA before the algorithm starting encrypting. This was done in order to remove any performance impact from external memory or serial communication links. The AXI Timer submodule kept a cycle-precise timer that was initialized directly before starting the encryption. After the decryption finished, the timer was stopped. This ensured that no serial communication time for reading out was included in the final time count.




Fig. 5. Execution Time per Implementation (seconds)




Fig. 7. Additional Resources per Implementation

C. Resource Utilization

Resource utilization sees a significant increase when going from the three stage pipeline to the five stage pipeline. Utilization increases about 33%, which correlates to the addition of 2 extra stages with extra logic and registers added with each stage. The benefit of Xilinx DSP slices can be seen from these results. Adding a multiplier in a pipeline sees virtually no increase in LUT or FF utilization, since the DSP slice performs the entire multiply. The additional benefit of these slices is that they have much smaller critical paths than if they were built out of traditional LUT or FF hardware. The dividers increase utilization by about 6%, but see a very large performance benefit compared to implementations without the divider. Finally, the addition of the BTB does not add significant logic, but does incur one block RAM in order to store the branch target.




Fig. 6. Resource Utilization per Implementation

D. Power Utilization

Power was compared from two different sources. The Xilinx provided power estimator built into the Vivado framework was used to get a estimation of power for each of the different tests.




Fig. 8. Vivado estimated Utilization per Implementation (mW)

In order to capture a more accurate representation, measurements were done during processing which allowed for a read out of the board power. In order to isolate the FPGA device power from these board measurements, a baseline power reading was done with no logic utilization. This established the baseline power draw from our system, allowing for an effective comparison between the estimated power, which only accounts for on-chip power.




Fig. 9. Measured Power Utilization per Implementation during runtime (mW)

As the ability to increase speed and logic increases, power becomes more of a concern for digital designs. In this example with one MicroBlaze on a relatively low end FPGA, these power numbers appear small. However, many systems will realize multi-core architectures or chip architectures with the potential to consume more power. The results shown from this study are important due to this scalability. Although there is significant overhead shown with static power and PLL power, this is with only 5% of the FPGA being used. With 100% utilization of the FPGA, the dynamic power will play a large role in determining the overall power. The less logic and slower the clock speed, the less power will be used. The 5 stage pipeline with BTB now demonstrates some advantage for certain systems. Due to its slower frequency, it was lower in performance even with the additional hardware. However, this also makes it consume less power, giving it a commendable performance per watt.




Fig. 10. Performance per Watt (kB/s per W)

VII. CONCLUSION

FPGA devices benefit greatly from the use of soft processors. These processors add greater system flexibility over a system without software components. Applications such as cryptography are often found on FPGAs in order to take advantage of the hardware acceleration alongside the traditional processor core. The critical component of these cores is what sacrifices will need to be made to achieve this flexibility. This concern is even greater in the embedded field, where power and speed are both costly. Clock frequency plays a large role in increasing performance and power utilization. Adding more hardware specific to an application's purpose, such as multipliers for cryptography, can strongly increase performance. However, a system's design cannot be guaranteed to see greater performance strictly by increasing the hardware; only hardware that directly targets the applications and instructions being processed in the core will see benefit. With a greater amount of logic fitting on a single reconfigurable device, these processor cores will continue to be a benefit to system designers who seek additional configurability and flexibility with their algorithms and designs.

ACKNOWLEDGMENT

For support during this undertaking, thanks to Dr. Mohamed El-Hadedy. Additional thanks to Md Mohsin for a starting point on RSA encryption implementation in C.

REFERENCES

- [1] S. M. Trimberger, "Three Ages of FPGAs: A Retrospective on the First Thirty Years of FPGA Technology," *Proceedings of the IEEE*, vol. 103, no. 3, pp. 318-331, Apr. 2015.
- [2] Henkel, J, "A low power hardware/software partitioning approach for core-based embedded systems," *Design Automation Conference (DAC)*, 1999.
- [3] Xilinx, "Zynq-7000 SoC First Generation Architecture," ds190 datasheet, July 2018.
- [4] R. Lysecky and F. Vahid, "A Study of the Speedups and Competitiveness of FPGA Soft Processor Cores using Dynamic Hardware/Software Partitioning," *Design, Automation and Test in Europe*, Mar. 2005.
- [5] Xilinx, "MicroBlaze Processor Reference Guide," *MicroBlaze Reference*, October 2017.
- [6] K. Vipin and S. A. Fahmy, "FPGA Dynamic and Partial Reconfiguration," *ACM Computing Surveys*, vol. 51, no. 4, pp. 1-39, 2018.
- [7] J. G. Tong, I. D. L. Anderson, and M. A. S. Khalid, "Soft-Core Processors for Embedded Systems," *2006 International Conference on Microelectronics*, Dec. 2006.